

Architecture

NALA Senior Analytics Engineer assessment. Karan Manoharan, September 2026.

The project transforms three sources (payments backend and fincrime service via StreamServe CDC, Amplitude via Fivetran) into a governed layer for four production requirements and one exploratory funnel. It runs unchanged on DuckDB (local, synthetic data) and Snowflake (production). Everything described here has been executed locally: 37 models, 1 snapshot, 158 data tests, 4 unit tests, semantic layer validated against the warehouse, five metrics queried.

Diagram: [docs/architecture.svg](#) in the repository.

1. Layering and materializations

Layer	Models	Materialization	Why
Staging (stg_)	16, one per source table	view	Rename, cast, collapse the CDC log. No business logic, so no storage.
Intermediate (int_)	4	view	Reusable logic (type policy, workflow context, task candidates). Only consumed by marts.
Snapshot	snap_transaction_state	snapshot (type 2)	Source keeps only the current row; this records state transitions.
Marts, facts (fct_)	5	incremental merge (transactions, attempts, rule executions); table (reviews, tasks)	Large tables get merge with a 3-day updated_at lookback. Small ones rebuild.
Marts, aggregates (agg_)	5	table	Cheap to rebuild from facts; one row per day and dimension for Hex.
Marts, dimensions (dim_)	3	table	Users, rule versions, providers.
Exploratory	3	table for the 90-day event subset, views on top	Cost bound on the 100M-row source; nothing downstream depends on it.

Production and exploratory differ in three places: the production selector never includes the Amplitude branch; exploratory tests are warn, production tests are error; contracts are enforced only on marts/. A failing exploratory model cannot block a finance number.

Grains are explicit and never mixed: transaction, disbursement attempt, rule execution, formal review, task, onboarding entity. Facts carry pre-aggregated 0/1 counters (completed_attempt_count, false_positive_count) so ratios are always sums divided by sums, never averages of rates.

2. CDC handling (StreamServe)

Assumption: StreamServe lands an append-only change log with _cdc_operation, _cdc_lsn, _cdc_loaded_at. Every staging model calls cdc_current_rows(): keep the row with the highest LSN per primary key, drop keys whose latest operation is DELETE. If the connector is configured to merge into a current-state table instead, cdc_landing_mode: current_state turns the macro into a pass-through. One variable, no model changes.

Consequences downstream:

- Source tests check not_null only; uniqueness is asserted on staging, where it is true.
- Incremental facts filter on updated_at, never created_at. An attempt created two months ago that changes state today falls inside the lookback and is merged.
- Hard deletes disappear from staging; the snapshot keeps their last known state.
- completed_at on transactions comes from the snapshot's first COMPLETED transition, falling back to updated_at when the current state is COMPLETED (updated_at is the last state change by definition).
- Freshness is measured on _cdc_loaded_at (warn 15 min, error 2 h) and _fivetran_synced (warn 26 h, error 30 h). Business timestamps are never used as freshness.

3. Cross-source enrichment

Requirement 4 joins fincrime workflow executions (Source 2) to backend tasks (Source 1). There is no foreign key. `tasks_task.associated_ids` is positional and only documented as "typically". The join is therefore a candidate match, not a lookup:

1. `int_fincrimetaskcontext` parses the typed positions (`transaction_review`: [0] transaction, [2] user; `user_review`: [0] user) and the `content.disbursement_id` anchor. Contradictory IDs are flagged.
2. `int_workflowcontext` extracts `user_id/transaction_id` from the execution JSON and whether the actions include `CREATE_TASK`, `BLOCK_USER` or `HOLD_TRANSACTION`.
3. `int_task_workflow_candidates` matches on the same entity and a task-producing action inside the 60 minutes before the task was created.
4. `fct_fincrimetasks` labels each task `UNIQUE_INFERRED`, `AMBIGUOUS`, `UNMATCHED`, `UNSUPPORTED_CONTEXT` or `CONFLICTING_CONTEXT`. All tasks stay in the count. Outcome-versus-result correlation is reported with linkage coverage beside it.

Challenges met: IDs live inside JSON on both sides (typed in staging, parsed once in intermediate); the two CDC streams have independent lag, so a task can land before its execution (the window tolerates it, the freshness gate catches large lag); positional arrays are a contract that can drift (a test fails on unsupported task types instead of guessing).

The fix is upstream: emit `workflow_execution_id` when the task is created. The model is built so that column replaces the heuristic in one join.

Amplitude joins to the backend through `user_id`. Anonymous events on a device seen with exactly one user are stitched for the POC; shared devices stay unresolved.

4. Orchestration

Scheduler: dbt Cloud jobs. The brief assumes dbt Cloud access, the sources are already managed by StreamServe and Fivetran, and there is no cross-tool dependency that needs an external orchestrator yet. If Fivetran and dbt need a shared DAG later, the project is already asset-shaped for Dagster.

Job	Selector	Cadence	Trigger	Notes
ops_hourly	tag:ops tag:fincrimetask tag:tasks + parents	hourly	schedule, skipped if the previous run is still going	incremental facts, 3-day lookback; snapshot first
finance_daily	tag:finance + parents	daily 01:00 UTC	schedule, after the FX rate for the previous day has landed (source freshness gate on <code>reference_data</code>)	full rebuild of finance aggregates; reconciliation tests block publication
growth_daily	selector:exploratory	daily	Fivetran sync-complete webhook, not a fixed 06:00	90-day window; failures warn only
full_refresh	everything	on demand	manual	after a policy seed change or a backfill

Dev to prod: feature branch, PR opens a dbt Cloud CI job that builds `state:modified+` into `PR_<n>` schema deferring unchanged models to the production manifest, runs tests, and tears the schema down on merge. GitHub Actions runs the same project on DuckDB (`make build, mf validate-configs, sqlfluff, dbt parse --target prod`) so a broken model never reaches the warehouse CI. Merge to `main` deploys; production jobs read `main`.

Cost controls: incremental merge on the three large facts, `cluster_by` date on Snowflake, exploratory subset materialized once per day, query tags per job, `store_failures` off.

5. Testing strategy

Layer	Minimum	Examples
Source	not_null on keys, freshness on connector timestamps, JSON parseability	<code>json_payloads_parse</code>

Layer	Minimum	Examples
Staging	unique + not_null on keys, accepted_values on every enum, relationships on foreign keys	17 state and type enums
Intermediate	grain uniqueness, no future or wrong-scope candidates	no_future_workflow_candidate
Production marts	all of the above plus reconciliation to the fact (finance_count_reconciliation, reviews_reconcile, fincrime_tasks_reconcile), denominator bounds, latency band sums, FX coverage, no partial USD totals, unit tests on the denominator rules	4 dbt unit tests, 23 singular tests
Exploratory	grain uniqueness, ordering, monotonic funnel, missing KYC signal; severity warn	kyc_final_signal_missing

A production mart ships when its grain test, its reconciliation test and its unit tests pass on Snowflake CI. An exploratory model ships when it builds and its grain test passes.

6. Semantic layer

Seven semantic models, all on facts or dimensions, never on aggregates: transactions, users, disbursement_attempts, rule_reviews, rule_execution_activity, fincrime_tasks, onboarding. Measures are the 0/1 counters and amounts; ratio metrics divide summed measures, so disbursement_provider_success_rate across providers is 91/110, not the mean of 90% and 10%.

Time axes are per measure: completed_transaction_volume aggregates on completed_date, transaction_success_rate on transaction_date, reviews on review_date. transactions carries a user foreign entity so any metric slices by sender_country from dim_users.

Saved queries finance_volume_by_corridor and provider_success_daily encode the safe groupings Hex should start from. Local-currency volume must always be grouped by sent_currency or corridor; MetricFlow cannot forbid a cross-currency sum, so the saved query and the docs do.

Assumptions

Topic	Assumption	How to confirm
CDC landing	Append-only log with _cdc_operation, _cdc_lsn, _cdc_loaded_at	StreamServe config; flip cdc_landing_mode if merged
Timestamps	UTC, stored as TIMESTAMP_NTZ	Postgres and connector settings
Outbound volume	DISBURSEMENT, CONVERSION_DISBURSEMENT, COLLECTION_CONVERSION_DISBURSEMENT, OUTGOING_PEER_TO_PEER, CONVERSION_OUTGOING_PEER_TO_PEER	Finance sign-off on P2P; policy is a seed
USD rates	RAW.REFERENCE.DAILY_FX_RATES (date, currency, usd_per_unit, source); latest rate on or before the day, max 7 days old	Finance owns the feed
Completion time	Snapshot history, else updated_at when state is COMPLETED	Add completed_at upstream if available
Attempt durations	last_updated_at - created_at for terminal states	State-event table if one exists
Task linkage	Positional associated_ids + 60-minute window	Emit workflow_execution_id on task creation
KYC complete	kyc_step.completed with event_properties.is_final_step = true	Product analytics contract

Topic	Assumption	How to confirm
Amplitude identity	Explicit user_id; single-user device stitched; shared unresolved	Growth team
Fivetran metadata	_fivetran_synced present	Always true for Fivetran